

## GreenBus PCI Express User's Guide

Copyright GreenSocs Ltd 2007

Developed by Christian Schroeder and Wolfgang Klingauf Technical University of Braunschweig, Dept. E.I.S.

20<sup>th</sup> November 2007



## **Contents**

1	Intr	oduction	3				
	1.1	Terms, Acronyms, Documentation and Implementation Conventions	3				
	1.2	PCI Express	3				
2	Gen	eral Project Handling	4				
	2.1	Global Project Settings	4				
	2.2	Include PCIe.h	5				
	2.3	Generic Devices in PCIe	6				
3	PCI	e specific Classes	7				
	3.1	Transactions and Accesses	7				
	3.2	Using the PCIe API	9				
		3.2.1 Processing incoming TLPs	9				
		3.2.2 Creating and sending TLPs	10				
	3.3	Configuration Space	12				
	3.4	Addressing	12				
	3.5	Switch					
	3.6	Root Complex					
	3.7	Top-Level testbench	14				
	3.8	3.8 Mix Generic and PCIe Devices					
		3.8.1 How to connect Generic Devices to PCIe Switches?	14				
		3.8.1.1 Introduction	14				
		3.8.1.2 User View	17				
		3.8.1.3 Alternative Connection	17				
		3.8.2 How to connect PCIe Devices to Generic Routers?	18				
	3.9	Specials	20				



3.9.1	Interrupts	20
3.9.2	Power Management Turn Off Message and Ack	20



## Chapter 1

## Introduction

This is the User's Guide of the PCI Express (PCIe) implementation for GreenBus. This documentation expects the reader to be familiar with SystemC and GreenBus. GreenBus is the universal bus implementation by GreenSocs.

The GreenBus PCIe implementation allows the user to model a PCIe topology at the programmer's view (PV) abstraction level.

The documentation is related to the GreenBus / PCIe version TODO.

# 1.1 Terms, Acronyms, Documentation and Implementation Conventions

PCIe PCI Express

TLP Transaction Layer Packet which is modeled with a PCIe Transaction

PCIe Transaction One PCIe TLP or several TLPs (correlated Request, Completions)

The implementation of all not-user classes according PCIe are located in the *namespace tlm::PCIe*.

## 1.2 PCI Express

This PCIe implementation is based on the PCI Express Base Specification Revision 2.0 and related documents.



## Chapter 2

## **General Project Handling**

This chapter is a short introduction how to prepare your project for PCIe.

## 2.1 Global Project Settings

For your global project settings we recommend to use a global setting file (e.g.  globals.h ) which is included by each header file before any other includes (expect systemc.h).
Debug Output
In the global setting file you can switch on the PCIe terminal outputs by defining
#define PCIeDEBUG_ON
To enable other GreenBus outputs use
#define VERBOSE

#### **Additional PCIe Checks**

The PCIe implementation is prepared to check additional rules apart from the obligatory ones defined in the PCIe specification. To enable these additional checks, use

#define CHECK\_RULES\_ENABLE



Note that these checks need some more run time. Disable these checks for large simulations where you already know that no failures occur.

### 2.2 Include PCIe.h

PCle.h:	∟ globals.h	and afterwards
<pre>#include "globals.h" #include "greenbus/transport/PCIe/PCIe.h"</pre>		

The second include defines the PCIe Transaction and access classes and (in the case of static casts) typedefs the generic types to PCIe.

Generic devices may include greenbus/transport/GP/GP.h instead of PCIe.h if some settings in globals.h are done, see section 2.3.



PCIe projects should be hard coded to use the PCIe Transactions					
by including PCle.h - see above. In the case of mixed projects					
with generic devices - which include GP.h - the user can en-					
able the PCIe Transaction with a make switch instead of including					
an user defined globals.h: make EXTENSION=PCIe . Also see					
greenbus/examples/README for more details.					

### 2.3 Generic Devices in PCIe

When combining generic and PCIe peripherals together within a PCIe project use the static casts version. In the generic peripherals' header files the user has to take care of using the PCIe Transaction definition either by global define or by make file switch (see section 2.2).

All generic types are typedefed to PCIe types (e.g. GenericTransaction = PCIeTransaction, GenericMasterAccess = PCIeMasterAccess).

In the case the dynamic casts version is used together with generic ports, routers, devices etc, define the following defines in the  $\square$  globals.h: #define EXTENDED\_TRANSACTION PCIe (Uses PCIe transaction), #define USE\_PCIE\_TRANSACTION (Enables typedefs Generic = PCIe)



## **Chapter 3**

## **PCIe specific Classes**

This chapter leads through the process of creating a project with PCIe Devices, Switches, Root Complex etc. In the end there is a section 3.8 about how to connect generic devices to PCIe and vice versa.

#### 3.1 Transactions and Accesses

Figure 3.1 shows the derivation tree for static casts PCIe Transactions and Accesses.

The PCIe types are typedefed to generic types:

```
typedef PCIe::PCIeTransaction GenericTransaction;
typedef PCIe::PCIeSlaveAccess GenericSlaveAccess;
typedef PCIe::PCIeMasterAccess GenericMasterAccess;
typedef PCIe::PCIeRouterAccess GenericRouterAccess;
typedef PCIe::PCIePhase GenericPhase;
```

The quark access methods in the PCIe Transaction class are protected to avoid the direct usage of the Transaction class. The user has to use the Access class which matches the actual usage type: When creating a transaction use the MasterAccess, when processing a received transaction, use the SlaveAccess. Smart pointers of these Access classes are typedefed to [Master|Slave]AccessHandle. Router Accesses are not covered in the document just as phases.

The user can use class methods to get Accesses out of a transaction:

```
PCIeMasterAccess &getMasterAccess();
PCIeSlaveAccess &getSlaveAccess();
```

And the user can use global methods to get Access handles from Transactions and Atoms. (Atoms are not used in the PV implementation):

```
GenericMasterAccessHandle _getMasterAccessHandle(/*Atom or Transaction*/)
GenericSlaveAccessHandle _getSlaveAccessHandle(/*Atom or Transaction*/)
```

How to use the access methods also see section 3.2.

In a PCIe peripheral module the user instantiates the PCIe API (see section 3.2).



## **Static Casts Derivation Tree**

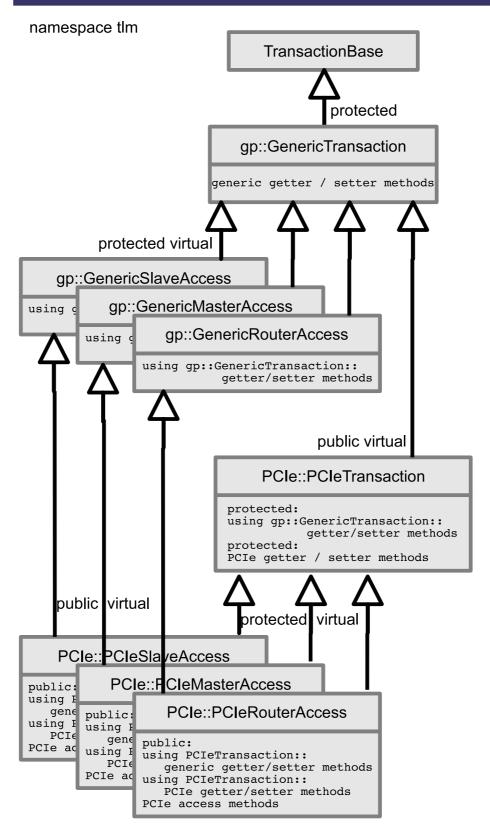


Figure 3.1: Derivation tree Transactions, Accesses.



When creating a PCIe TLP the PCIe Master Access class provides convenience methods for all existing TLP types, e.g.:

```
init_Memory_Read(/* address, data, data size, requester func. no. */);
init_Memory_Write(/* address, data, data size, requester func. no. */);
init_IO_Read(/* address, data, data size, requester function no. */);
init_IO_Write(/* address, data, data size, requester function no. */);
init_Configuration_Read(/* bus no., dev. no., func. no., ..., data, ...*/)
;
init_Configuration_Write(/* bus no., dev. no., func. no., ..., data, ...*/)
);
init_Message(/* message code, requester func. no. */);
```

When receiving and processing a PCIe TLP the PCIe Slave Access class provides convenience methods, e.g.:

```
get_TLP_type();
get_addr();
init_Memory_Read_Completion(/* completion status, completer func. no. */);
// ...
set_Completion_Status(/* status */)
init_Unsupported_Request();
```

For the detailed and complete list of access methods look at the doxygen API reference.

### 3.2 Using the PCIe API

The PCIe framework provides one API (PCIeAPI) for PCIe peripherals independent from their main behavior as a master or slave since each PCIe device is a master *and* a slave in (Green)Bus terms. Include greenbus/api/PCIe/PCIeAPI.h before using the API.

## Implementation Note Bidirectional PCIe Ports

Because each PCIe device is a bus master and a slave the PCIe API has one bidirectional port which acts as a master *and* a slave. The bidirectional port (PCIeBidirectionalPort) is a GenericInitiatorAPI which is instantiated with the constructor parameter is\_bidirectional=true.

### 3.2.1 Processing incoming TLPs

The PCIe API creates a PCIe Configuration Space Header of type 0 (for PCIe device Functions) and handles the incoming Configuration Request TLPs automatically. The user module needs not to care about them. Details about Configuration Space see section 3.3.

The API does some additional checks for incoming transactions. Some of the rule checks can be switched off by removing the define PRECHECK\_ENABLE in \_\_\_\_ greenbus/api/PCle/PCleAPI.h . Af-



ter these first checks the transaction is forwarded to the <code>b\_transact</code> method which has to be implemented by the user module. The user device has to implement the <code>PCIe\_recv\_if</code> interface and give the this pointer to the API during its construction. Use the convenience methods of the Slave Access class to process the TLP in the switch statement. An example implementation of a sending and receiving PCIe device is presented in the example listing in figure 3.2. See an example for a receiving device e.g. in <code>greenbus/examples/PCIe/platform/PCIeRecvinfoDevice.h</code>.

#### 3.2.2 Creating and sending TLPs

The same PCIe API can be used for creating and sending TLPs.

Call the <code>create\_transaction()</code>-method on the API to get a Master Access Handle to the TLP which should be sent. Afterward call the convenience method(s) according to the TLP type you want to send. Call <code>myAPI.send(myMasterAccessHandle)</code> on the API to send the TLP. In a final step process the completion of the TLP if there is one again using the convenience methods of the Access class.

See figure 3.2 for an example listing. Find an example for a sending device e.g. in greenbus/examples/PCle/platform/PCleSendDevice.h.

If a device should only send but not process received transactions the constructor of the API can be called with NULL instead of the this pointer. Then the API will internally process the transactions (e.g. handle Configuration Requests and manipulate the Configuration Space) but will not try to forward them to the user device implementation. If a transaction is received which needed user processing an sc\_warning will be reported.



```
|#include "my_globals.h"
 #include "greenbus/transport/PCIe/PCIe.h"
 #include "greenbus/api/PCIe/PCIeAPI.h"
 using namespace tlm;
s using namespace tlm::PCIe;
 class MyPCIeDevice
  : public sc_module,
    public tlm::PCIe::PCIe_recv_if // for receiving ability {
10 public:
   PCIeAPI myAPI;
   SC_HAS_PROCESS(MyPCIeDevice);
   MyPCIeDevice(sc_module_name name)
   : sc_module(name),
      myAPI("PCIeAPI", this) // for sending and receiving
     SC_THREAD(send_action); // for sending
   /// Method for receiving ability
   virtual void b transact(PCIeTransactionHandle th) {
     PCIeSlaveAccessHandle ah = _getSlaveAccessHandle(th);
     switch((unsigned int)th->get_TLP_type()) {
       case MemWrite:
       // ...
   /// Thread for sending ability
   void PCIeSendDevice::main_action() {
     // dummy data
     std::vector<gs_uint8> *dat;
     dat = new std::vector < gs_uint8 > (data_size);
     for (unsigned int i = 0; i < data_size; i++) dat->at(i) = i;
     // create transaction
     PCIeMasterAccessHandle ah;
     ah = mAPI.create transaction();
     // fill transaction with data
     ah->init_IO_Write( addr, *dat, data_size );
     // send transaction
     mAPI.send transaction(ah);
     // process Completion
     if (ah->get_Completion_Status() == SuccessfulCompletion) {
       /* process data */
     } else { /* Completion not successfull */ }
     delete dat; dat = NULL;
```

Figure 3.2: Example listing PCIe device.



### 3.3 Configuration Space

TODO: The Configuration Space is not yet implemented. We are waiting for DRF.

### 3.4 Addressing

The PCIe framework supports the three address (routing) types of PCIe: address based, ID based and implicit routing.

The PCIe Transaction has two boolean variables which store the kind of used address type due to fast access (figure 3.3):

	Address based Routing	ID based routing	Implicit routing
mIsIDbasedRouting	false	true	false
mIsImplicitRouting	false	false	true

Figure 3.3: Routing types.

#### Address based routing

The address based routing is divided into two independent address spaces: *Memory address space* and *I/O address space*. The routing information is transported in the quark mAddr which is the default quark for addresses in GreenBus. GreenBus routing with MAddr matches address based routing. Other routing mechanisms cannot be mapped to MAddr field routing because all 64 bit are needed by address based routing.

#### **ID** based routing

The PCIe ID based routing is mapped to the quarks mBusNo, mDevNo, mFuncNo and mRegNo. ID based routing is needed for Configuration Requests. The Requester ID is the ID of a TLP's sender. The user needs not to handle these addresses. Only the Root Complex implementation gets in contact with the IDs. Requester IDs are set automatically by the API.

#### Implicit routing

Implicit routing is done with the quark mMessageType. The user needs not to handle implicit routing. The message creating convenience methods set the correct values.



RoutedToRootComplex	000	Route to Root Complex
RoutedByAddress	001	Never used
		(mIsIDbasedRouting=mIsImplicitRouting=false)
RoutedByID	010	Route normally by ID
		(set mIsIDbasedRouting=true
		instead of mIsImplicitRouting!)
BroadcastFromRootComplex	011	Broadcast from Root Complex to all devices
LocalTerminateAtReceiver	100	Local - Not forwarded at the next switch
GatheredAndRoutedToRootComplex	101	see PME_TO_Ack Message Code:
		Collect Acks from Downstream Ports
		and send one to Upstream Port

Figure 3.4: Implicit routing types.

#### 3.5 Switch

The PCIe Switch class is named PCIeRouter following the GreenBus terms. A Switch has an Upstream Port (upstream\_port) which has to be connected to another bidirectional PCIe port: a Downstream Port of another Switch or a Root Complex. The Downstream Port of a Switch is a multi port. The user can connect one or more bidirectional PCIe ports to this port: mainly that will be PCIe device's or other Switch's Upstream Ports. (The Downstream port is not needed to be connected to any port. It may left unbound.)

Generic devices may be connected to the Downstream Port according the rules in section 3.8.

### 3.6 Root Complex

The Root Complex is a device with an internal PCIe Switch. To help the user with implementing the Root Complex there is a class PCIeRootComplex (see \_\_\_\_greenbus/api/PCIe/PCIeRootComplex.h ). This derives from the PCIeRouter class, accordingly the user is able to connect devices and switches to its Downstream Port. Do not connect the Upstream Port: it is bound internally to the down\_to\_router\_port which is another PCIeAPI port. That port is the connection from the software/CPU to PCIe and should be used by the user to implement the Root Complex functionality.

At minimum the user has to implement the method  $PCIeRootComplex:: Q down_to_router_port_b_transact(PCIeTransactionHandle th). To be able to extend the Root Complex the user should derive from this class, e.g. MyPCIeRootComplex: public PCIeRootComplex. See greenbus/examples/PCIe/platform for an example.$ 

- The PCIeRootComplex is a (derives from) PCIeRouter.
- Use Downstream Port to connect the PCIe tree topology.
- Do not use the Upstream Port which is connected internally.
- Use the port down\_to\_router\_port to access the PCIe topology.



- Implement down\_to\_router\_port\_b\_transact to handle incoming TLPs from the PCIe topology.
- Derive from PCIeRootComplex to extend functionality, e.g. to add a SC\_METHOD.

### 3.7 Top-Level testbench

The PCIe topology is built in the top-level testbench.

Rules:

- Create exactly one Root Complex instance.
- Create Switches if needed.
- Connect each Switch's Upstream Port to another Switch (Downstream Port) or to the Root Complex (Downstream Port).
- Create PCIe peripherals (devices) and connect their Upstream Ports either to a Switch's or a Root Complex' Downstream Port.

Figure 3.6 shows an PCIe example testbench with connected Root Complex, Switches and devices. The example can be found in the platform example greenbus/examples/PCIe/platform and is illustrated in figure 3.5.

#### 3.8 Mix Generic and PCIe Devices

This section describes how a generic device can be connected to a PCIe topology using the extension mechanism.

#### 3.8.1 How to connect Generic Devices to PCIe Switches?

The main precondition for getting generic devices compatible to be connected to a PCIe topology is that the generic device has to use PCIe Transactions and ports.

#### 3.8.1.1 Introduction

In a generic device the included header is GP.h (see section 2.2). This include has to typedef PCIe Transactions (and other classes) to generic ones. This is done either by changing the include to Greenbus/transport/PCIe/PCIe.h or by using the make switch (see section 2.2). Afterwards PCIe transactions are created instead of generic transactions while keeping the same generic API.

Generic transactions are address based so the compatibility API creates PCIe MemoryRead/Write transactions in parallel to generic transactions.



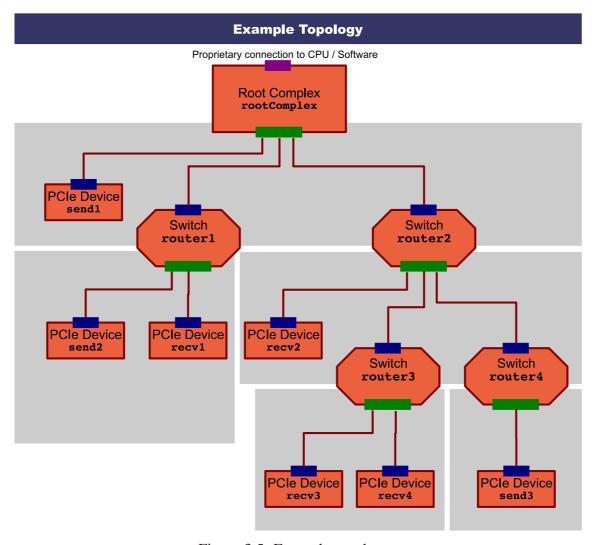


Figure 3.5: Example topology.

#### Implementation Note

#### Generic Devices in a PCIe topology

For this behavior, the convenience access methods in the PCIe Access classes which are used by the generic device have to be adapted. E.g. setMCmd(..) sets the PCIe quark MPCIeCmd instead of (more precisely in addition to) MCmd.

Now the generic devices can be connected to the PCIe port of the router. If the generic device has two ports (initiator\_port and target\_port) it has to be connected to the router twice. This is no issue for configuration because generic devices don't understand any ID routed TLPs anyway.

Generic Read and Write Transactions have an equivalent in PCIe: Memory Read and Write Transactions. Due to this analogy it is possible to map generic Generic\_MCMD\_RD and Generic\_MCMD\_WR transactions to PCIe MemRead and MemWrite transactions.



```
#include "my_globals.h"
 #include "greenbus/transport/PCIe/PCIe.h"
 #include "greenbus/transport/PCIe/PCIeRouter.h"
 #include "MyPCIeRootComplex.h"
#include "PCIeSendDevice.h"
 #include "PCIeRecvDevice.h"
 int sc_main(int, char**)
   // instantiate partizipants
10
   MyPCIeRootComplex rootComplex("rootComplex");
   PCIeRecvInfoDevice recv1("recv1");
   PCIeRecvInfoDevice recv2("recv2");
   PCIeRecvInfoDevice recv3("recv3");
   PCIeRecvInfoDevice recv4("recv4");
   PCIeSendDevice send1("send1");
   PCIeSendDevice
                     send2("send2");
   PCIeSendDevice
                     send3("send3");
   PCIeRouter router1("router1");
   PCIeRouter router2("router2");
   PCIeRouter router3("router3");
   PCIeRouter router4("router4");
   // connect Devices
   rootComplex.downstream_port(send1.mAPI);
   router1.upstream_port(rootComplex.downstream_port);
   router1.downstream_port(send2.mAPI);
   router1.downstream_port(recv1.mAPI);
   router2.upstream_port(rootComplex.downstream_port);
   router2.downstream_port(recv2.mAPI);
   router3.upstream_port(router2.downstream_port);
   router3.downstream port(recv3.mAPI);
   router3.downstream_port(recv4.mAPI);
   router4.upstream_port(router2.downstream_port);
   router4.downstream_port(send3.mAPI);
   // start simulation
   sc_start();
   return 0;
```

Figure 3.6: Example top-level testbench.



#### Implementation Note

#### Mapping Generic Command Type to PCIe Command type

PCIe uses an own Command enumeration because the generic one is not extendable. The two analogical transactions have the same number.

The supported features are: Read, Write Transactions, Error handling, Completion / Bytes Valid.

#### **3.8.1.2** User View

To connect modules which use generic ports (initiator\_port, target\_port,...) to PCIe modules or routers the user has to do the following steps (see example === examples/PCIe/generic/):

- Make sure to use the PCIe typedefs for generic classes in the generic modules (either include PCIe.h or use make switch).
- Instantiate the modules and routers as usual in the testbench.
- Manually set the address range of the generic module:

```
gen_s.target_port.base_addr = (MAddr) 0x00000000000000000001L;
gen_s.target_port.high_addr = (MAddr) 0x000000000000FFFFLL;
```

Connect the generic device to a PCIe Router Downstream Port like it is done with PCIe modules, e.g.

```
router.downstream_port(gen_m.init_port);
router.downstream_port(gen_s.target_port);
```

■ Generate ID map after "all" connections are done:

```
router.m_PCIeAddressMap.generate_ID_map();
```

■ Manually register the generic address range at the router:

```
router.add_address_range(gen_s.target_port);
```

#### 3.8.1.3 Alternative Connection

An alternative to the direct connection is the PCIe2GenericPortWrapper which was developed for connecting a PCIe Device to a Generic Router (see section 3.8.2). The wrapper maps the two ports of a master and slave Generic Device to one bidirectional PCIe port.

Figure 3.7 shows the available alternatives connecting a Generic Device to a PCIe Router.



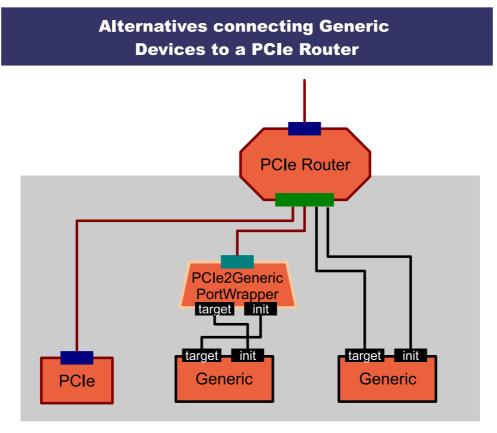


Figure 3.7: Alternative connections Generic Device to PCIe Switch

#### 3.8.2 How to connect PCIe Devices to Generic Routers?

A PCIe Device can be connected to a generic topology (e.g. router) if the generic modules use PCIe classes instead of generic ones (same precondition as in section 3.8.1). The PCIe Device only may send Memory Read and Memory Write TLPs. At the PCIe Device received generic reads and writes are mapped to PCIe Memory Reads and Memory Writes.

If the PCIe device is only a master or a slave (only sends or receives TLPs) the device may be connected directly to the <code>init\_port</code> or <code>target\_port</code> of the Generic Router. The user has to set the <code>base\_addr</code> and <code>high\_addr</code> parameters of the API port manually.

If the PCIe device should act as master *and* slave a wrapper (PCIe2GenericPortWrapper) has to be used to map the bidirectional PCIe port to an init\_port and a target\_port. Now the target\_port of the wrapper automatically gets the address range of the PCIe port (during before\_end\_of\_elaboration).

This wrapper class in in greenbus/protocol/PCle/PCleGenericWrapper.h . The listing in figure 3.8 shows how to use this wrapper.



```
#include "greenbus/protocol/PCIe/PCIeGenericWrapper.h"

// ... see greenbus/examples/PCIe/generic/

tlm::PCIe::PCIe2GenericPortWrapper wrapper_pcie_m(")
    PCIe2GenericPortWrapper_pcie_m");
PCIeSendDevice3 pcie_m("PCIe_Master");

// Set Device's address range
    pcie_m.mAPI.base_addr = (MAddr) 0x000000000100000LL;
pcie_m.mAPI.high_addr = (MAddr) 0x000000000100000LL;

// Connect PCIe Device with wrapper (as Master and Slave)
    r.init_port(wrapper_pcie_m.target_port);
    wrapper_pcie_m.init_port(r.target_port);
    wrapper_pcie_m.pcie_port(pcie_m.mAPI);
```

Figure 3.8: Connect PCIe Device to Generic Router with Wrapper.

An example using this wrapper (and comments for not using the wrapper) is in the example greenbus/examples/PCle/generic/ (define USE\_GENERIC\_ROUTER\_CONNECT).

Figure 3.9 shows the available alternatives connecting a PCIe Device to a Generic Router.

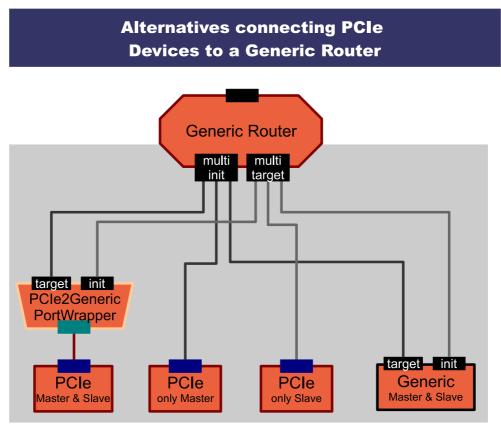


Figure 3.9: Alternative connections PCIe Device to Generic Router



## 3.9 Specials

This sections is a collection of some special TLPs etc.

#### 3.9.1 Interrupts

Legacy support interrupts are provided by with INTx Messages (Assert/Deassert) which can be generated with the transaction access methods. The PCIe Router has to handle these Interrupt messages in a special way.

The PCIe Router supports a basic mechanism handling INTx Assert and Deassert Messages: An (Downstream) incoming Interrupt Message is terminated at the receiver (Switch) (see MessageType). The Interrupt is forwarded after calculating the mapping. This approach does not match all requirements (see Specification pages 72f.) but will care for forwarding the correct Message. (Not supported e.g.: Interrupt Disable bit of the Command register, Deassert Message in the case of DL\_Down, Root Complex support.)

#### 3.9.2 Power Management Turn Off Message and Ack

The acknowledgment to the Power Management message PME\_Turn\_Off is a special routing case: The Turn\_Off message is broadcasted to all devices and each device has to react with a PME\_TO\_Ack message. A Switch has to forward the Ack only after having received all Acks from its downstream devices. A Switch counts the incoming PME\_TO\_Ack messages and sends the PME\_TO\_Ack message to its upstream port after having received Acks from all Downstream Ports.